

Question 1

Let $\phi : X \times A \rightarrow \mathbb{R}^d$

The task is to compute $\nabla_{\theta} \log(\pi_{\theta}(a|x))$ for the Boltzmann policy:

$$\pi_{\theta}(a|x) = \frac{e^{\theta^{\top} \phi(x,a)}}{\sum_b e^{\theta^{\top} \phi(x,b)}} \quad (1)$$

To answer this, we take logs and the derivative:

$$\nabla_{\theta} \log \pi_{\theta}(a|x) = \nabla_{\theta} \log \frac{e^{\theta^{\top} \phi(x,a)}}{\sum_b e^{\theta^{\top} \phi(x,b)}} \quad (2)$$

$$= \nabla_{\theta} \log e^{\theta^{\top} \phi(x,a)} - \nabla_{\theta} \log \sum_b e^{\theta^{\top} \phi(x,b)} \quad (3)$$

$$= \phi(x,a) - \frac{1}{\sum_b e^{\theta^{\top} \phi(x,b)}} \sum_{b'} e^{\theta^{\top} \phi(x,b')} \phi(x,b') \quad (4)$$

$$= \phi(x,a) - \sum_{b'} \frac{e^{\theta^{\top} \phi(x,b')}}{\sum_b e^{\theta^{\top} \phi(x,b)}} \phi(x,b') \quad (5)$$

Finally obtaining:

$$\nabla_{\theta} \log(\pi_{\theta}(a|x)) = \phi(x,a) - \sum_b \pi_{\theta}(b|x) \phi(x,b) \quad (6)$$

Which can be rewritten as:

$$\nabla_{\theta} \log \pi_{\theta}(a|x) = \phi(x,a) - \mathbb{E}_{\pi_{\theta}}[\phi(x, \cdot)] \quad (7)$$

Question 2

For this question, we built a policy gradient algorithm with a Boltzmann policy using $\phi_i(x, a) = 1_{(a=1)}$. In this section, we introduce the algorithm. We then answer the questions posed in the exercise in the following sections.

The MDP we study has one state x (that is, $X = \{x\}$) and two actions $\mathcal{A} = \{a_1, a_2\}$. The rewards are defined as follows:

$$R_k = \begin{cases} 1, & \text{with probability } p_{I_k} \\ 0, & \text{otherwise} \end{cases}$$

The rewards are received with probability $p_1 = 1/2$ and $p_2 = 1/2 + \Delta$. Given the environment set-up, we define a learning algorithm in order to find the optimal policy π_{θ_k} . We use the policy gradient method to learn the optimal weight parameters $\theta_{k+1} = \theta_k + \alpha_k g_k$. The arguments our algorithm takes are:

1. Enviroment: Takes delta as argument and react to actions paying 1 or 0 with probability $p_1 = 1/2$ and $p_2 = 1/2 + \Delta$.
2. Gamma: Discount factor for future rewards.

3. Policy: In this case we used Boltzmann policy as defined previously
4. Seed: We choose a random seed to make sure it is possible to reproduce the simulations.
5. Step type: Here we defined a fixed learning rate and also different rate decay in function of the number of steps.
6. Delta: The environment parameter that gives how much the reward from action 2 is better than action 1.
7. Weights θ : Weights of our policy
8. Beta: Learning rate for the value function

The following plots show some training results obtained by averaging over various training runs, using different random states to initialize.

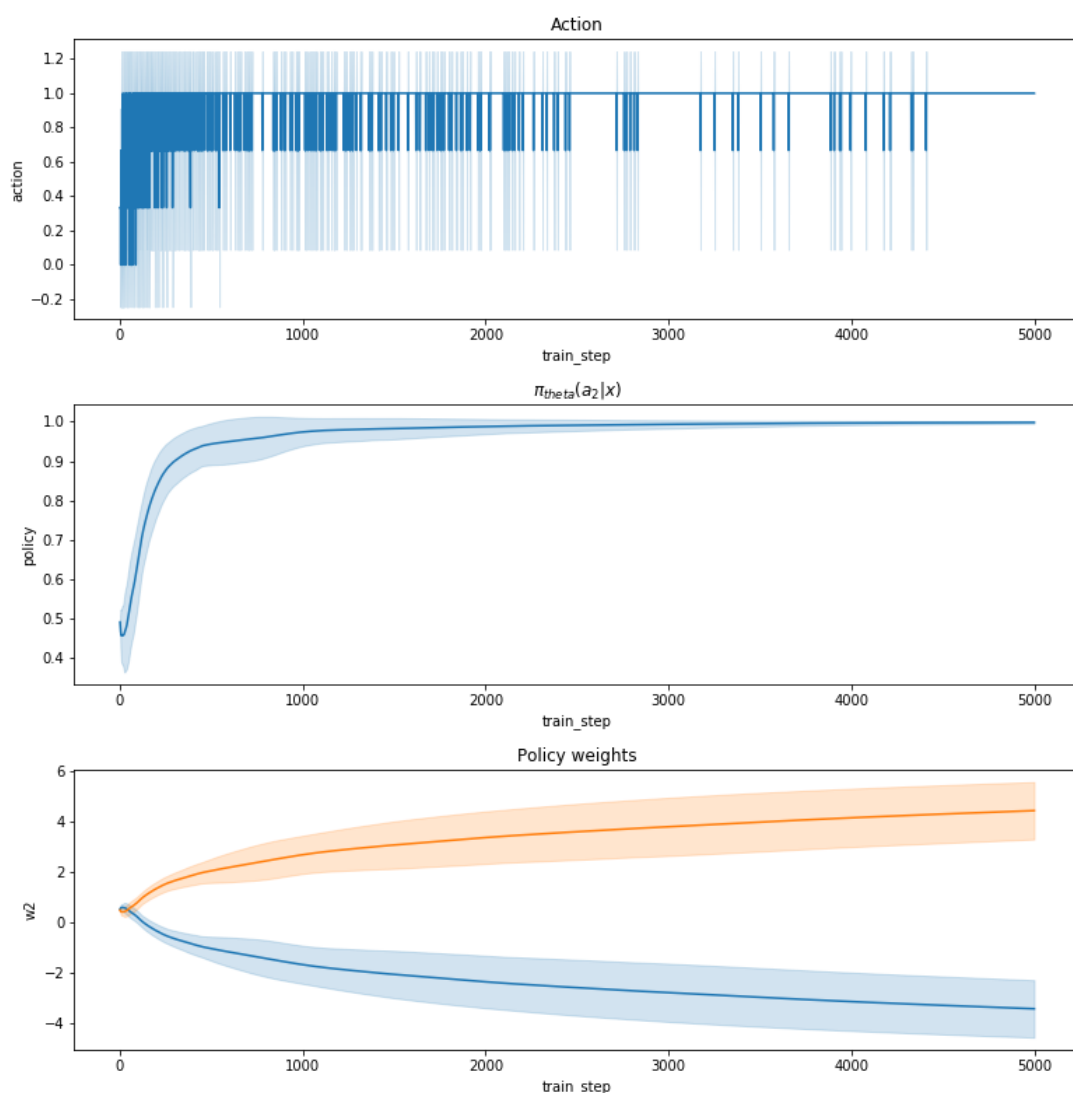


Figure 1: Convergence of the algorithm

The first plot shows the mean of actions taken over the training steps and the corresponding standard deviation. The second plot shows the convergence of the policy function and the

third shows the evolution of the weights θ .

Note that the policy values are plotted together with the standard deviation over different seeds. Policy function values cannot be higher than 1 or lower than 0, however, mean \pm sd can be higher than 1. The main objective is to illustrate the volatility of the measure.

Having introduced the algorithm we built, we now move to answer the questions for the exercise.

Question 2.1

Task: What is the influence of γ on the algorithm.

We find that γ does not play a role in this policy gradient. We get this result because we use the particularly simple estimation of the expectation in the policy gradient function suggested in the hint. When we set $h(x) = \gamma V^{\pi_{\theta_k}}(x)$, take a single sample to approximate the expectation, and take $\widehat{Q}^{\pi_{\theta_k}}(x, A_k) = R_k + \gamma V^{\pi_{\theta_k}}(x)$, the expression for the policy gradient simplifies as follows:

$$\nabla_{\theta} \rho(\theta) = \sum_{(x,a) \in \mathcal{X} \times \mathcal{A}} \mu_{\theta}(x) \pi_{\theta}(a|x) \nabla_{\theta} \log \pi_{\theta}(a|x) (Q^{\pi_{\theta}}(x, a) - h(x)) \quad (8)$$

$$= \mathbb{E}_{X \sim \mu_{\theta}, A \sim \pi_{\theta}(\cdot|X)} [\nabla_{\theta} \log \pi_{\theta}(A|X) (Q^{\pi_{\theta}}(X, A) - h(X))] \quad (9)$$

$$\approx \nabla_{\theta} \log \pi_{\theta}(a|x) (\widehat{Q^{\pi_{\theta}}(x, a)} - h(x)) \quad (10)$$

$$= \nabla_{\theta} \log \pi_{\theta}(a|x) (R_k + \gamma V^{\pi_{\theta_k}}(x) - \gamma V^{\pi_{\theta_k}}(x)) \quad (11)$$

$$= \nabla_{\theta} \log \pi_{\theta}(a|x) R_k \quad (12)$$

$$= \left(\phi(x, a) - \sum_b \pi_{\theta}(b|x) \phi(x, b) \right) R_k \quad (13)$$

The discount factor γ cancels out. Further, since the state in this MDP does not change, we do not have to take into account the consequences of our actions now for our future as we adjust our policy. Therefore the choice of γ should intuitively not play a role in our policy gradient.

Question 2.2

Task: Fix $\gamma = 0.99$ and $\delta = 0.05$, and consider the step sizes $\frac{\alpha_k}{\sqrt{k}}$, $\frac{\alpha_k}{k}$, and $\frac{\alpha_k}{k^2}$ for various choices of c . Plot $\pi_{\theta_k}(a_2|x)$ as a function of k .

We find that, for small delta, only the fixed step size manages to get the right answer 100% of the times. Using the other step sizes some seeds converge to the wrong policy, for a given initial c . The issue here is that the learning rate decreases too fast to reach the optimal parameters, so we get stuck in the beginning. On the other hand, for larger learning rates, the policy jumps in the beginning and then gets stuck in a suboptimal level as well.

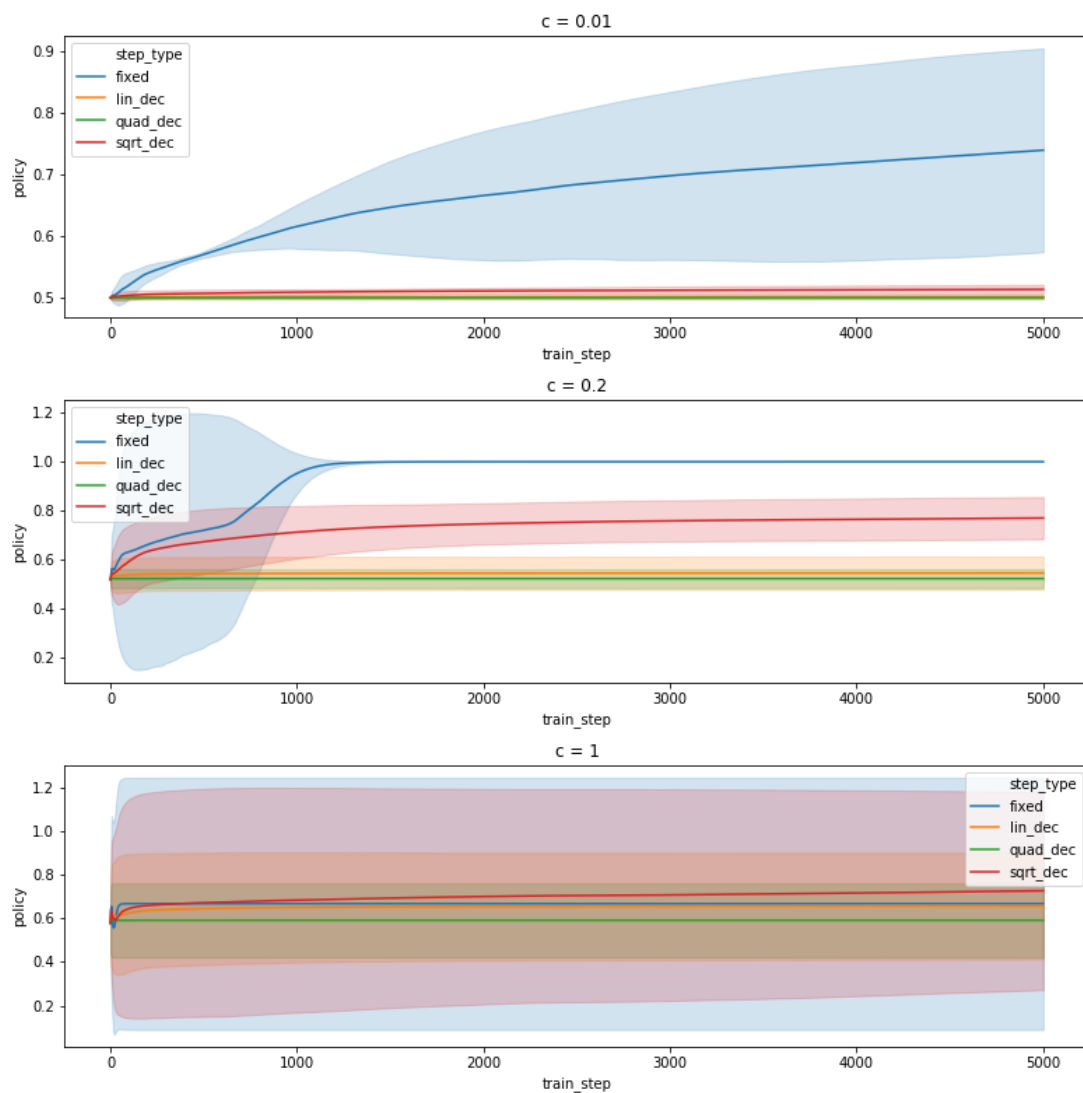


Figure 2: Influence of the step size

The optimal learning rate also depends on the parameters we are estimating. Since we are starting with small θ , large learning rates send the parameters around by too much.

Question 2.3

Task: Find the best step size and policy as function of Δ .

We observe faster convergence for big deltas. This makes intuitive sense, since the "signal" from choosing the correct action is stronger for large deltas on average. For lower deltas we see that we sometimes don't even converge to the optimum.

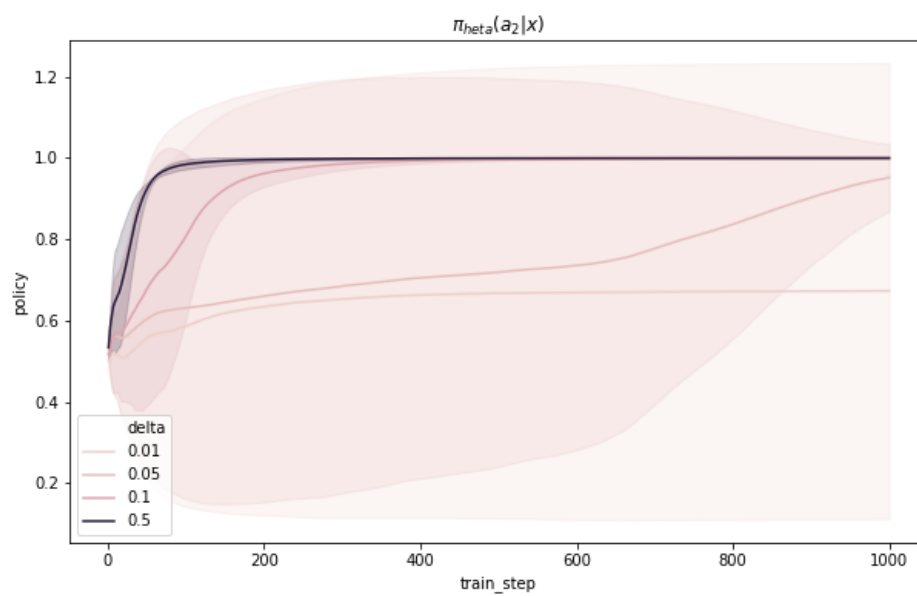


Figure 3: Influence of delta

Question 3

Task: Performance of standard UCB and ϵ -greedy algorithms

Question 3.1

First we define the standard UCB algorithm for our bandit problem. As described in the question, at each iteration of the algorithm, we start by choosing an "arm" or action in this case by finding

$$I_t = \arg \max_{i \in \{1,2\}} \left(\hat{\mu}_{t,i} + \sqrt{\frac{2 \ln t}{N_{t,i}}} \right)$$

we initialize the algorithm by forcing two trials of each action. After taking an action we store the resulting reward and update our mean estimator $\mu_{t,i}$.

The following plot shows the total reward for $T = 10,000$ rounds as a function Δ of the environment:

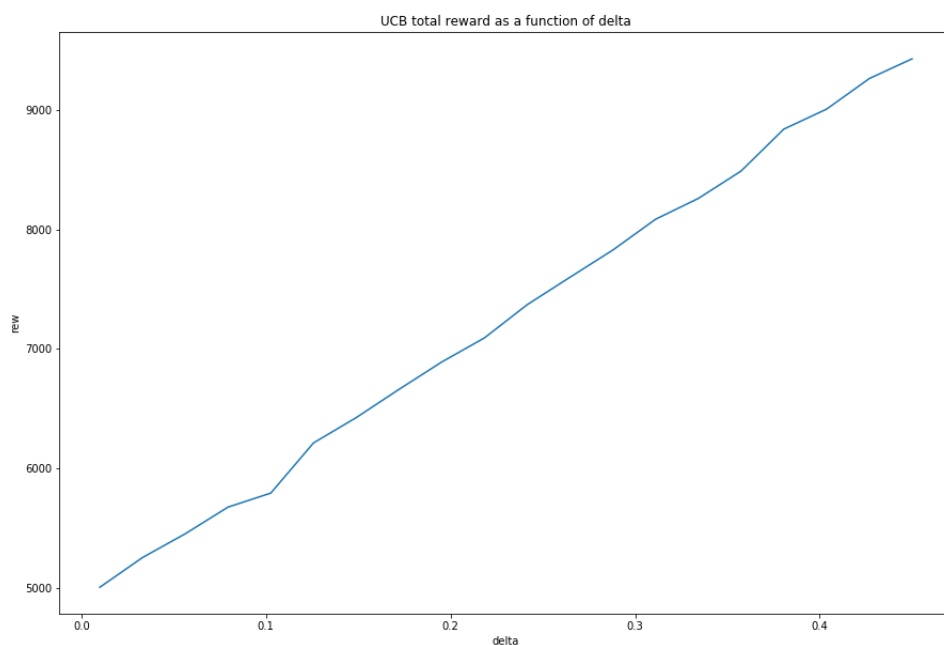


Figure 4

One can easily see that larger values of Δ lead to greater total rewards. The intuition behind this is that large Δ lead to a clear separation between the expected rewards of each action. For very small Δ , the expected rewards of each action are almost $\frac{1}{2}$. Unsurprisingly, the total reward for a Δ of 0 is 5000. Conversely, for a Δ of 0.5 the total reward is almost 10,000, but not quite. This is due to the fact that we still have to run a tiny amount of trials to get convergence of the estimated mean to the true value. After that we always pick the right action, which almost certainly results in a reward.

Question 3.2

The ϵ -greedy algorithm works in the following way:

For a given probability ϵ , either choose a uniformly random action, or choose the action which has the maximum expected reward based on the previously recorded rewards. The value of ϵ drops in each step, resulting in a smaller and smaller chance of taking random actions. Consequently, the drop in ϵ can be seen as a switch from exploring the environment to exploiting it based on the reward history. ϵ in this case is set as $\epsilon = c/t$ where c is a constant.

The following plot shows the total reward for a range of values of c and Δ .



Figure 5

First we note that similarly to the UCB algorithm, larger values of Δ lead to higher total rewards. Additionally, for small c and Δ all total rewards are similar. For greater Δ , with smaller parameter c , the algorithm converges quicker to the optimal policy. This can be seen in the higher overall rewards. Again, the intuition is that for high Δ we don't need to explore the environment as in depth as before, so smaller c will lead to a smaller amount of trials in which we randomly explore the environment.

Question 3.3

The last plot shows the combined results of all three algorithms for our small bandit problem with $T = 10000$ for all runs.




Figure 6

All algorithms have similar total rewards for a given Δ with UCB and Policy-Gradients being slightly better than ϵ -greedy for larger Δ . This is probably due to the fact that the problem at hand is a simple one, with only one state and two possible actions. In this case we don't lose too much if we use a "simple" policy like UCB, over the more complex policy gradient method. Further, finding the optimal learning rate for the policy gradient method is hard, whereas UCB is much simpler to implement and fine tune.